

**What you should already know**

In this lesson, we will build on what you have already learned about the Python programming language. To be sure that we are all on the same page, let's briefly review the things about Python that you should already be familiar with. For more detail, review the lesson on *Introduction to Computer Programming*.

**Data types, constants, and variables**

You should know that the kinds of values that can be expressed in a programming language are known as its **data types**. The **primitive types** of a programming language are those data types that are built-in (or standard) to the language and typically considered as basic building blocks (i.e., more complex types can be created from these primitive types). Python's standard types can be grouped into several classes: numeric types, sequences, sets, and mappings. You should be familiar with numeric types and sequences (e.g., lists).

You should know that a **constant** is defined as a value of a particular type that does not change over time. In Python both numbers and text may be expressed as constants. **Numeric constants** are composed of the digits 0 through 9 and, optionally, a negative sign (for negative numbers), and a decimal point (for floating point numbers). **Text constants** consists of a sequence of characters (also known as a string of characters – or just a **string**).

You should know that a **variable** is defined to be a named object that can store a value of a particular type. Before a variable can be used, its name must be declared.

**Input and output**

You should be familiar with obtaining input (via the `input` function) and generating output (via the `print` statement) in Python. Here's a simple example:

```
name = input("What is your name? ")
print("Hello, {}".format(name))
```

**Operators**

You should be familiar with a variety of operators in Python. Specifically, arithmetic operators, relational (comparison) operators, and assignment operators. Arithmetic operators include addition, subtraction, etc, and perform arithmetic operations on operands. Relational operators include comparison of equality, inequality, less-than, and so on, and perform comparisons on operands and return true or false. Assignment operators include operators such as `+=`, `-=`, and so on, and combine assignment with arithmetic.

**Primary control constructs**

You should be very familiar with the three primary control constructs: sequence, selection, and repetition. Sequence implies one statement after another. Selection allows blocks of optional statements to be executed. Repetition provides a mechanism for repeating blocks of statements. There are two main forms of repetition that we have covered: iteration and recursion. Iteration involves repeating a task some fixed number of times, until a condition is reached, or over some structure (such as the items in a list). Although recursion was only briefly covered, you should know that it involves breaking a

problem down repeatedly into smaller versions of itself until a base or trivial case is reached. We will cover recursion in much more detail later in the curriculum.

## Subprograms

You should be quite familiar with subprograms and how they can encapsulate behavior in programs. They are organized, reusable, and related statements that perform some action. Specifically, some subprograms perform tasks and terminate; others return a value. You should understand how control flow is transferred to a subprogram when a subprogram is called, and how it is returned when the subprogram terminates.

## A review of program flow

Although you should be familiar with this already, it is so important that we should probably go over it in detail again. It is very important to be able to identify the flow of control in any program, particularly to understand what is going on. In fact, this significantly helps to debug problems in programs. Recall that, in Python, function definitions aren't executed in the order that they are written in the source code. Functions are only executed when they are called. This is perhaps best illustrated with an example that you have seen before:

```
1:  def min(a, b):
2:      if (a < b):
3:          return a
4:      else:
5:          return b

6:  def max(a, b):
7:      if (a > b):
8:          return a
9:      else:
10:         return b

11: num1 = int(input("Enter a number: "))
12: num2 = int(input("Enter another number: "))
13: print("The smaller is {}".format(min(num1, num2)))
14: print("The larger is {}".format(max(num1, num2)))
```

Each Python statement is numbered for reference. Lines 1 through 5 represent the definition of the function `min`. This function returns the *minimum* of two values provided as parameters. Lines 6 through 10 represent the definition of the function `max`. This function returns the *maximum* of two values provided as parameters. Lines 11 through 14 represent the main part of the program. Although the Python interpreter does *see* lines 1 through 10, those lines are not actually *executed* until the functions `min` and `max` are actually called. The first line of the program to actually be executed is line 11. In fact, here is the order of the statements executed in this program if `num1 = 34` and `num2 = 55`:

11, 12, 13, 1, 2, 3, 14, 6, 7, 9, 10

Let's explain. Line 11 asks the user to provide some value for the first number (which is stored in the variable `num1`). Line 12 asks the user to provide some value for the second number (which is stored in the variable `num2`). Line 13 displays some text; however, part of the text must be obtained by first calling the function `min`. This transfers control to line 1 (where `min` is defined). The two actual

parameters, *num1* and *num2*, are then passed in and mapped to the formal parameters defined in `min`, *a* and *b*. Then, line 2 is executed and performs a comparison of the two numbers. Since *a* = 34 and *b* = 55, then the condition in the if-statement is true. Therefore, line 3 is executed before control is transferred back to the main program with the value of the smaller number returned (and then control continues on to line 14). Note that lines 4 and 5 are never executed in this case!

Line 14 is then executed and displays some text. Again, part of the text must be obtained by first calling the function `max`. This transfers control to line 6 (where `max` is defined). The variables *a* and *b* take on the values 34 and 55 respectively. Line 7 is then executed, and the result of the comparison is false. Therefore, line 8 is not executed. Control then goes to line 9, and then to line 10 which returns the larger value. The program then ends.

What is the order of execution if *num1* = 55 and *num2* = 34?

What if *num1* = 100 and *num2* = 100?

Here's another example with a simple for loop:

```
1: for a in range(1, 4):
2:     for b in range(1, 5):
3:         print("{} * {} = {}".format(a, b, a * b))
```

This snippet of code displays a portion of a multiplication table. In fact, here's the output:

```
1 * 1 = 1
1 * 2 = 2
1 * 3 = 3
1 * 4 = 4
2 * 1 = 2
2 * 2 = 4
2 * 3 = 6
2 * 4 = 8
3 * 1 = 3
3 * 2 = 6
3 * 3 = 9
3 * 4 = 12
```

Here's the order of the statements executed. To make things a bit more clear, it is grouped and highlighted:

1, 2, 3, 2, 3, 2, 3, 2, 3, 1, 2, 3, 2, 3, 2, 3, 2, 3, 1, 2, 3, 2, 3, 2, 3, 2, 3  
 |----- 1 \* n -----|----- 2 \* n -----|----- 3 \* n -----|

The first portion (highlighted in red and labeled  $1 * n$ ) represents a single iteration of the **outer** for loop and a *full* iteration of the **inner** for loop. It generates the following output:

```
1 * 1 = 1
1 * 2 = 2
1 * 3 = 3
1 * 4 = 4
```

Let's explain. Line 1 of the outer for loop generates the list  $[1, 2, 3]$ . It then iterates over the values in the list with the variable  $a$  taking on each value, one at a time. Initially,  $a = 1$ . Line 2 represents the inner for loop and generates the list  $[1, 2, 3, 4]$ , and iterates over its values with the variable  $b$  taking on each value, one at a time. Initially,  $b = 1$ . Line 3 then displays the first line of output:  $1 * 1 = 1$ . This makes sense because  $a$  and  $b$  are both 1.

So far, the order of statements executed is 1, 2, 3. Note that the inner for loop iterates its complete cycle (i.e., through the entire generated list) for each iteration of the outer for loop. Therefore, the inner for loop iterates through the list  $[1, 2, 3, 4]$  for each value in the outer for loop's list  $[1, 2, 3]$ . After line 3 (when  $a$  and  $b$  are both 1), the inner for loop then iterates to the next value in the list. Therefore, line 2 is executed again so that  $b = 2$ . Similarly, line 3 executes again, generating the output:  $1 * 2 = 2$ . At this point, the order of statements executed is 1, 2, 3, 2, 3.

The inner for loop continues iterating two more times (lines 2 and 3), setting  $b$  to 3 and then to 4. After the first full iteration of the outer for loop, the order of statements executed is 1, 2, 3, 2, 3, 2, 3, 2, 3. So when  $a = 1$ ,  $b$  goes through the values 1, 2, 3, and 4. The output generated at this point is then:

```
1 * 1 = 1
1 * 2 = 2
1 * 3 = 3
1 * 4 = 4
```

Since the inner for loop has finished a full iteration, control goes back to line 1, thereby allowing the outer for loop to iterate to the next value in the list so that  $a = 2$ . Line 2 is executed again, generating a new list  $[1, 2, 3, 4]$  and setting  $b$  to 1. Similarly, line 3 is executed again, generating the output:  $2 * 1 = 2$ . Lines 2 and 3 are executed as before, for each value in the inner loop's list  $[1, 2, 3, 4]$ . Of course, this generates exactly the same order of statements as before: 1, 2, 3, 2, 3, 2, 3, 2, 3. However, this represents the second iteration of the outer loop.

Since the inner for loop has finished another full iteration, control goes back to line 1, thereby allowing the outer for loop to iterate to the next value in the list so that  $a = 3$ . Again, lines 2 and 3 are executed as before, for each value in the inner loop's list  $[1, 2, 3, 4]$ . Clearly, this generates exactly the same order of statements as before: 1, 2, 3, 2, 3, 2, 3, 2, 3. This time, it represents the third (and final) iteration of the outer loop. Why? Because the outer for loop has iterated through the entire list  $[1, 2, 3]$ . After this final iteration, the variable  $a$  has taken on all of these values. Therefore, the outer for loop is exited, and the program terminates.

Again, knowing the order in which statements are executed is crucial to debugging programs and ultimately to creating programs that work.

This concludes a review of what you should already know in Python. From here, we'll introduce new content.

## Formal vs actual parameters

You have seen that a function can have parameters. These parameters are formally stated when the function is defined; for example:

```
def average(a, b):  
    return (a + b) / 2.0
```

Here, the variables *a* and *b* are formally defined as parameters that must be passed in to the function `average` when it is called. In this context, the variables *a* and *b* are called **formal parameters**. It is where they are defined (in a formal manner).

**Definition:** *Formal parameters are formally defined as part of the function header.*

Now consider a point in the source code where this function is called; for example:

```
avg = average(11, 67)
```

Here, the result of a call to the function `average` with the supplied values (or parameters) 11 and 67 is stored in the variable *avg*. These values, 11 and 67, are considered **actual parameters** in this context. That is, they are the *actual* values that will be passed in as parameters to the function `average`.

**Definition:** *Actual parameters (or arguments) are the actual values that are passed into a function.*

In fact, they are mapped to the formally defined parameters (i.e., formal parameters) *a* and *b* in the function `average`. That function will use these values to make calculations and return the average of the two. The value returned replaces the function call. Think of this replacement as follows:

```
avg = average(11, 67)  
      39.0
```

Therefore, the variable *avg* is assigned the value 39.0 after the call to the function `average` is complete. Consider this call to the same function:

```
x = 11  
y = 67  
avg = average(x, y)
```

Here, the result is still the same. The average of the two variables, *x* and *y* (with the values 11 and 67 respectively), is stored in the variable *avg*. Here, *x* and *y* are also actual parameters (even if they are variables themselves) because they represent the actual values supplied to the function `average`.

## Variable scope

Consider the following Python program snippet:

```
a = 10  
  
def f(x):  
    a = 11  
    b = 21  
    x *= 2
```

```

    print("in f(): a={}, b={}, x={}".format(a, b, x))

b = 20
f(b)
print("in main: a={}, b={}".format(a, b))

def g():
    global a
    a *= 1.5
    print("in g(): a={}, b={}".format(a, b))

g()
print("in main: a={}, b={}".format(a, b))

```

The variables *a* and *b* (highlighted above) are considered **global variables**.

**Definition:** *Global variables* are accessible throughout the entire program because they are defined outside of any block context (e.g., a loop construct, a function, etc). Global variables can be accessed anywhere. Their **scope** is global (i.e., throughout the entire program).

**Definition:** A variable's **scope** is the region of code that variable is accessible in.

Take a look at the output of the program above:

```

in f(): a=11, b=21, x=40
in main: a=10, b=20
in g(): a=15.0, b=20
in main: a=15.0, b=20

```

Let's explain the output. Initially, the variable *a* is assigned the value 10. The next segment of code defines the function *f*. This is only a definition (i.e., the statements are not actually interpreted or executed at this point). Then, the variable *b* is assigned the value 20. What follows is a call to the function *f*, passing the variable *b* as an **actual parameter**. Control is then transferred to the function *f*, whose statements are now executed. Note that, to the function *f*, the variable *x* is the formal parameter that takes on the value passed in (from the variable *b*). So the variable *x* is now equal to the value of the variable *b* (i.e., 20) that was passed in at the point of the call to *f*. Note that the variable *x* is local to the function *f*; therefore, it is considered a **local variable**. That is, it is defined in *f* and only accessible in *f* – its **scope** is valid only in the function *f*.

**Definition:** A **local variable** is a variable defined in (and therefore only accessible in) a local region/context. For example, a variable defined inside a function is **local** to that function and therefore only accessible within that function.

Also note that, although *a* and *b* are global, there are **local** versions declared in *f*. It is important to note that these are different variables than the global versions – even if they have the same name!

**Definition: Shadowing** is the process of introducing a local variable that shares the same name as a global variable. This local variable “shadows” (or hides) access to the global variable of the same name. We will see this concept appear again later in the course, but with a different context.

So what happens in `f`? The local variable `a` is initialized with the value 11, the local variable `b` is initialized with the value 21, and the local variable `x` (which is passed in as an argument with the value 20) is doubled to 40. The output of the function `f` is then clear:

```
in f(): a=11, b=21, x=40
```

Once `f` completes and control is transferred back to the point at which function `f` was called, the variable `x` is no longer accessible! In fact, let's alter the print statement immediately after the call to `f` from:

```
print("in main: a={}, b={}".format(a, b))
```

And change it to:

```
print("in main: a={}, b={}, x={}".format(a, b, x))
```

Here's the output of the program now:

```
in f(): a=11, b=21, x=40
Traceback (most recent call last):
  File "scope.py", line 11, in <module>
    print("in main: a={}, b={}, x={}".format(a, b, x))
NameError: name 'x' is not defined
```

Notice the error indicating that the variable `x` is not defined. That's because it was defined in `f`; however, the current context is outside of `f`. The variable `x` is no longer available once `f` finishes and control is transferred back to the main part of the program.

Let's replace the print statement to remove the error and explain the rest of the output from the original execution of the program:

```
in main: a=10, b=20
in g(): a=15.0, b=20
in main: a=15.0, b=20
```

Once control is transferred back to the main part of the program, the local variables `a` and `b` (in `f`) no longer exist. However, the global variables `a` and `b` do! They were initialized to 10 and 20 respectively. Therefore, the next line of output makes sense:

```
in main: a=10, b=20
```

The next part of the program defines another function, `g`, that is then called. Note the **global** keyword in the function `g`. This instructs Python to reference a globally defined version of the variable that follows the global keyword. That is, a local version is not defined and/or initialized. Instead, the global version is directly referenced. Moreover (and quite importantly), it permits the global version to be **changed**. Although there are no arguments to the function `g`, the global variable `a` is directly **modifiable** through the global keyword. When the statement `a *= 1.5` is executed, the value of the global variable `a` is 10. This statement changes its value to 15.0, directly updating the variable's value – globally!

Note the print statement in `g`. It refers both to the variables `a` and `b`. A reference to the variable `a` makes sense; however, the variable `b` is also accessible. In fact, the variable `b` is referencing the global version of `b`, similar to the variable `a` (i.e., it is directly **readable**). The difference in using the global keyword is that it permits a change to the variable; without it, it can only be utilized in a read-only manner. Since the global variable `b` is initialized with the value 20, the output in `g` is clear:

```
in g(): a=15.0, b=20
```

When control is transferred back to the main part of the program, changes to global variable `a` persist (even if they were changed in a function!):

```
in main: a=15.0, b=20
```

To illustrate this even more, let's slightly change the function `g` as follows:

```
def g():
    global a
    a *= 1.5
    b = 40
    print("in g(): a={}, b={}".format(a, b))
```

Note the slight difference: `b` is declared and initialized with the value 40. Which `b` is this? Is it a local version (i.e., local to `g`)? Or is it referring to the global version declared in the main part of the program? Recall that, without the global keyword, a global variable can not be modified. Therefore, an assignment statement in a function to a variable that has the same name as a global variable indicates that the variable is a new instance, defined **locally** in the function. This is a different variable `b`! The output in `g` is clear:

```
in g(): a=15.0, b=40
```

When control is transferred back to the main part of the program, the local version of `b` disappears. All that's left is the global version (that remains unchanged at 20). Therefore, the output in the main part of the program is also clear:

```
in main: a=15.0, b=20
```

## A review of the Python list

Although you should be familiar with Python lists, they are quite important and used often; therefore, we will go over it again. Generally, a Python **sequence** is composed of (typically related) elements. Each element in a sequence is assigned an index (or position). A sequence with  $n$  elements has indexes 0 to  $n-1$ . Python has many built-in types of sequences; however, the most popular is called the list.

The **list** in Python is quite versatile. Recall that a list is declared using square brackets; for example:

```
grades = [ 94, 78, 100, 86 ]
```

The statement above declares the list `grades` with four integers: 94, 78, 100, and 86. The list can be displayed in its entirety (e.g., with the statement `print(grades)`); however, we can access each element individually by its index (specified within brackets). Accessing can mean to read a value in the list, or it can mean to change a value in the list; for example:

```
print(grades[0])
grades[3] = 87
```



```
grades[1] += 2
```

Recall that more than one value in a list can be accessed at a time. We can specify a range (or interval) of indexes in the format `[lower:upper+1]` which means the interval `[lower, upper)` (i.e., closed at lower and open at upper). That is, the lower index in the range is inclusive but the upper is not. For example:

```
stuff[3:4]      # accesses index 3 (the same as stuff[3])
stuff[0:5]      # accesses indexes 0 through 4
stuff[-3]       # accesses the third index from the right
```

Also recall that list elements can be deleted with the **del** keyword as follows:

```
del stuff[2]
```

Finally, recall that Python provides several built-in operations that can be performed on lists. Here are many of them:

<code>len(list)</code>	Returns the length of a list
<code>max(list)</code>	Returns the item in the list with the maximum value
<code>min(list)</code>	Returns the item in the list with the minimum value
<code>list.append(item)</code>	Inserts item at the end of the list
<code>list.count(item)</code>	Returns the number of times an item appears in the list
<code>list.index(item)</code>	Returns the index of the first occurrence of item
<code>list.insert(index, item)</code>	Inserts an item at the specified index in the list
<code>list.remove(item)</code>	Removes the first occurrence of item from the list
<code>list.reverse()</code>	Reverses the items in the list
<code>list.sort()</code>	Sorts a list

### Revisiting searching and sorting in Python

In previous lessons, we designed several searching algorithms (sequential/linear search and binary search) and sorting algorithms (bubble sort, selection sort, and insertion sort). We first specified them in pseudocode, and for some we showed how they could be implemented in Python (sequential search, binary search, and selection sort). To help get a better understanding of Python, let's briefly revisit some of these.

First, here's the **sequential search** for the smallest value in a list (from an earlier lesson). Note that a Python list is first populated with 20 random integers (from 1 to 99, also from an earlier lesson):

```
1: from random import randint

2: numbers = []
3: while (len(numbers) < 20):
4:     numbers.append(randint(1, 99))

5: print(numbers)
```

```

6:  minIndex = 0

7:  for index in range(1, len(numbers)):
8:      if (numbers[index] < numbers[minIndex]):
9:          minIndex = index

10: print("The smallest value is at index: {}".format(minIndex))
11: print("The smallest value is: {}".format(numbers[minIndex]))

```

This version of the sequential search technically returns the **index** of the smallest value (which is typically what programmers are interested in). Since the value can be easily accessed through the index, returning the index is much more meaningful. To generalize the sequential search so that it can return the index of a *specified* value (as opposed to the smallest value), it can be modified by replacing lines 6 through 11 as follows:

```

num = int(input("What integer would you like to search for? "))

for index in range(len(numbers)):
    if (numbers[index] == num):
        print("The value {} was found at index {}".format(num, index))

```

What happens if the specified value is duplicated several times in the list? Clearly, each index would be displayed. Here's example (with user input highlighted in **red**):

```

[20, 47, 80, 52, 98, 80, 1, 14, 31, 48, 70, 31, 97, 30, 31, 43, 59, 2, 38, 50]
What integer would you like to search for? 31
The value 31 was found at index 8!
The value 31 was found at index 11!
The value 31 was found at index 14!

```

But what if it's only necessary to find the *first* occurrence of a specified value (and then abort)? Python provides a way to **exit a repetition construct early** (a.k.a. an “early exit”) through the **break** keyword! Formally, the break keyword exits the *nearest enclosing* repetition construct. More on this in a bit. To illustrate the use of the break keyword, the sequential search code above can be modified to return only the first instance of a specified value:

```

for index in range(len(numbers)):
    if (numbers[index] == num):
        print("The value {} was found at index {}".format(num, index))
        break

```

Here's an example:

```

[87, 44, 37, 69, 92, 74, 49, 97, 65, 69, 27, 61, 22, 77, 3, 3, 25, 86, 53, 45]
What integer would you like to search for? 3
The value 3 was found at index 14!

```

Note that the value 3 occurs twice in the list (at index 14 and index 15); however, only the first instance is reported to the user before the search terminates. The break statement exits the enclosing repetition construct: in this case, the for loop.

What if the break keyword is located in a repetition construct that is also located inside of another repetition construct? In this case, it will exit the inner repetition construct only. Here's an example:

```

for i in range(2):
    print("i={}".format(i)),
    for j in range(5):
        print("j={}".format(j)),
        if (j > 1):
            break
    print()

```

Here's the output:

```

i=0 j=0 j=1 j=2
i=1 j=0 j=1 j=2

```

The outer for loop iterates *i* from 0 through 1. The inner for loop iterates *j* from 0 through 4. Moreover, the inner for loop exits early if *j* is greater than 1. Technically, the print statement in the inner for loop will display values of *j* that are less than or equal to 1. So why is a value of 2 for *j* displayed? When *j* is 2, the value is displayed, after which the if statement is executed (which breaks out of the inner for loop). The outer for loop continues (the lone print() statement is there to add a line break in between increasing values of *i*), and *i* becomes 1. This occurs again until the outer for loop terminates (when *i* is 2).

Let's now take a look at the binary search that was also covered in an earlier lesson. Recall that it is a very efficient search that requires a list to be sorted. Here's the Python code that was developed in an earlier lesson:

```

num = int(input("What integer would you like to search for? "))
found = False
first = 0
last = len(numbers) - 1

while (first <= last and found != True):
    mid = (first + last) // 2
    if (num == numbers[mid]):
        found = True
    elif (num > numbers[mid]):
        first = mid + 1
    else:
        last = mid - 1
if (found):
    print("{} was found at index {}".format(num, mid))
else:
    print("{} was not found.".format(num))

```

This version of the binary search keeps tracks of two boundaries (*first* and *last*) that identify the beginning and end indexes of the current portion of the list. Initially, *first* is 0 and *last* is *n*-1 (i.e., the entire list). If the middle value of the current portion of the list does not match the specified value, the appropriate half of the list is “discarded” by modifying either *first* (to discard the left half) or *last* (to discard the right half).

Recall that the binary search required a list to be sorted, thereby taking advantage of the algorithm's efficiency improvement over the sequential search. Here's the **selection sort** that was developed in an earlier lesson:

```
n = len(numbers)
for i in range(0, n - 1):
    minPosition = i
    for j in range(i + 1, n):
        if (numbers[j] < numbers[minPosition]):
            minPosition = j
    temp = numbers[i]
    numbers[i] = numbers[minPosition]
    numbers[minPosition] = temp
```

Recall how the selection sort works: (1) the list is sorted from left to right; (2) at each pass (controlled by the outer for loop), the smallest value is swapped with the first item in the unsorted portion of the list; and (3) the inner for loop performs the comparison of every remaining item in the unsorted portion of the list to find the smallest value. For a review, see the lesson on *Searching and Sorting*.

There were two other sorting algorithms that were covered in earlier lessons: bubble sort and insertion sort. We never developed Python code for them. Let's do this now. First, the bubble sort. Here is a version in pseudocode:

```
for i ← 1..list length-1
    for j ← 1..list length-i
        if item j of list < item j-1 of list
            then
                temp ← item j of list
                item j of list ← item j-1 of list
                item j-1 of list ← temp
            end
    next
next
```

You may not have seen a for loop described in pseudocode before; however, this is a common way to accomplish this repetition construct in pseudocode. So what's happening here? The basic idea is that a value in the list will be compared to the one before it. If they are out of order, then they are swapped. This continues, one index over (to the right), until the end of the list is reached. After the first pass, the largest value is guaranteed to be in its final position (i.e., at the end of the list). The next pass starts again at the beginning of the list; however, this time comparisons and swaps only take place until the second-to-last value in the list (because the last value has already been placed there during the last pass). Each time, the sorted list grows from right-to-left until the entire list is sorted.

The outer for loop controls the number of passes, while also providing a way to reduce the size of the unsorted portion of the list after each pass. It iterates from 1 through  $n-1$ . The inner loop controls the comparisons and swaps. Initially, the inner loop begins at 1 (the index of the second value in the list), and compares this value to the one before it (the first value in the list). If they are out of order, they are swapped. The swap works by using a variable (*temp*) that temporarily takes on one of the values in the list. This continues with the next index (i.e., 2), and so on. The last index compared is  $n-i$ . If the algorithm is in the first pass (i.e.,  $i$  is 1) and the length of the list is 10, the last valid index in the list is  $10 - 1 = 9$ .

Let's take a look at what a Python version of the bubble sort looks like:

```
n = len(list)
for i in range(1, n):
    for j in range(1, n - i + 1):
        if (list[j] < list[j - 1]):
            temp = list[j]
            list[j] = list[j - 1]
            list[j - 1] = temp
```

Recall that Python's range function uses the first parameter as a lower bound and the second parameter as one above the upper bound. That is, it operates on the interval  $[a, b)$ , where  $a$  is the (closed) lower bound and  $b$  is the (open) upper bound. Therefore, the upper bound of the outer loop is  $n$ : it iterates from 1 through (and including)  $n-1$  as intended. Similarly, the upper bound of the inner loop is  $n-i+1$ : it iterates from 1 through (and including)  $n-i$  as intended. In the inner loop, if any value at an index is less than the value of the one before it, they are swapped.

Next, let's take a look at the Python code for the insertion sort. Recall that the insertion sort works somewhat as you would arrange a hand of cards being dealt to you: a new card is inserted in its appropriate position in the hand of cards dealt so far.

```
1: i = 1
2: while (i < n):
3:     if (list[i - 1] > list[i]):
4:         temp = list[i]
5:         j = i - 1
6:         while (j >= 0 and list[j] > temp):
7:             list[j + 1] = list[j]
8:             j -= 1
9:         list[j + 1] = temp
10:    i += 1
```

Here's an explanation of the code. Line 2 controls the number of passes through the list ( $n-1$  total passes). The variable  $i$  is initialized to 1 (the second index in the list) and iterates through (but not including)  $n$ ; therefore, through the last index in the list. So, starting with the second value in the list, it looks to the left (of this current value). Line 3 checks if that value is greater, and if so, then it must move it to the right. Line 4 temporarily stores the current item, and lines 5 and 6 then begin the process of iterating from the previous element, continuing to the left. At any point, if a value to the left is greater than the current item, it is shifted one index to the right. This continues either until (1) the beginning of the list is reached; or (2) a value that is not greater is found. Ultimately, the current item is placed into its proper position in the list. The outer loop then continues with the next value in the list (through the last value in the list).

Note that there are many other ways that the searches and sorts shown could have been implemented in Python. For example, the use of for loops in the selection sort could have been replaced with while loops (or vice versa in the insertion sort).

### Other operators

Python provides several more classes of operators than you are already familiar with. Recall that, so far, you have been exposed to (and should be quite familiar with) arithmetic operators, relational

(comparison) operators, and assignment operators. In this lesson, we will cover several other classes of operators: **logical operators** and **membership operators**.

**Definition:** The **logical operators** evaluate two operands and return the logical result (i.e., True or False).

Think back to the primitive logic gates (*and*, *or*, and *not*). It turns out that they can be effectively mapped to conditions in if-statements. Logical operators operate on conditions (that use relational operators) and provide the overall logical result. In the following table, assume that  $a = \text{True}$  and  $b = \text{False}$ :

Python Logical Operators and Examples		
and	logical <i>and</i>	a and b is False
or	logical <i>or</i>	a or b is True
not	logical <i>not</i>	not a is False; not b is True

Note that this is equivalent to the primitive logic gates, where 0 is substituted for False and 1 for True. Here is the truth table for the and gate shown in this manner:

A	B	A and B
False	False	False
False	True	False
True	False	False
True	True	True

The logical operators sometimes make more sense when they are used in the context of a condition (e.g., in an if statement). Suppose that  $a = 5$  and  $b = 10$ . The following condition would evaluate to False:

```
if (a == 5 and b < 10):  
    ...
```

Why? Clearly because, although  $a$  is equal to 5,  $b$  is not less than 10 (it's equal to 10). Therefore, the *and* logical operator will return False if and only if both sides of the operator evaluate to True. In this case, the left side does while the right side does not. However, the following condition would evaluate to True:

```
if (a == 5 or b < 10):  
    ...
```

The *or* logical operator will return True if either (or both) sides of the operator evaluate to True. Since  $a$  is equal to 5, then the left side is True. In this case, the right side doesn't need to be evaluated (and, in fact, it isn't – more on that below).

The logical operators do work when the inputs (i.e.,  $a$  and  $b$  in the previous examples) aren't necessarily equal to True and False. That is, they also work when they are numeric values. Take, for example, the following statements:

```
1: a = 23  
2: b = 13
```

```

3:  print(a and b)
4:  print(a or b)
5:  print(not a)
6:  print(not b)
7:  a = 0
8:  print(not a)

```

Here's the output (with lines numbers matching those of the print statements above):

```

3:  13
4:  23
5:  False
6:  False
8:  True

```

The output of lines 3 and 4 can be a bit confusing. Why, for example, is *a and b* 13? Or why is *a or b* 23? This can be explained by the following table, where the variables *a* and *b* have numeric values (as in the examples above):

Python Logical Operators and Examples		
and	logical <i>and</i>	returns a if a is False, b otherwise
or	logical <i>or</i>	returns b if a is False, a otherwise
not	logical <i>not</i>	returns False if a is True, True otherwise

The output of lines 5, 6, and 8 makes sense when we realize that, in Python, 0 is False and any non-zero value is True! When *a* is 23 and *b* is 13, *a* evaluates to True (since it is non-zero); therefore, **not** *a* evaluates to False. This is the same with *b*. However, when *a* is 0, then it evaluates to False; therefore, **not** *a* evaluates to True. Formally, in the context of Boolean expressions, the following values are interpreted as false: False, None, numeric zero of all types, and empty strings and containers. All other values are interpreted as true.

### Did you know?

The *and* and *or* logical operators are **short circuit** operators. That is, to evaluate a True or False result, the minimum number of inputs required to produce such an output is evaluated. For example, suppose that *a* = False and *b* = True. The expression *a and b* is only True if both *a* and *b* are True. Since *a* is False, then there is no need to evaluate (or test) the value of *b*. This would be useless and waste CPU cycles. Similarly, if *a* = True and *b* = True, the evaluation of the expression *a or b* only requires checking that *a* is True for the entire expression to evaluate to True (i.e., there is no need to evaluate/test the value of *b*).

**Definition:** *Membership operators* test for some value's membership in a sequence (e.g., to test if an element exists in a list, or if a character exists in a string).

In the following table, suppose that the Python list *numbers* = [ 1, 3, 5, 7, 9 ], *x* = 2, and *y* = 3.

### Python Membership Operators and Examples

in	Returns True if a specified value is in a specified sequence or False otherwise	x in numbers is False; y in numbers is True
not in	Returns True if a specified value is not in a specified sequence or False otherwise	x not in numbers is True; y not in numbers is False

You have seen this in previous for loop examples (e.g., `for i in list`). This for loop configuration has the variable *i* take on each of the values in *list*, one at a time.

### String methods

Strings are often necessary when writing programs. As such, Python provides a variety of methods that work on strings. You have already seen one such method, **format()**, that formats a string as specified (we did this earlier in one variant of the **print** statement). The following table lists some of the more useful string methods:

Python String Methods/Functions	
<code>str.capitalize()</code>	capitalizes the first character of a string
<code>str.find()</code>	returns the first index of a string within another string
<code>str.format()</code>	formats a string according to a specification
<code>str.isdigit()</code>	determines if a string consists only of numeric characters
<code>str.lower()</code>	converts a string to lowercase
<code>str.replace()</code>	replaces all occurrences of a string (within a string) with another string
<code>str.split()</code>	returns a list of the words in a string
<code>str.upper()</code>	converts a string to uppercase

These string methods are explained in greater detail in a variety of online sources. We suggest that you Google them and try them out. However, here are a few examples in IDLE:

```
>>> s = "So, when is this going to get difficult?"
>>> s
'So, when is this going to get difficult?'
>>> s.capitalize()
'So, when is this going to get difficult?'
>>> s.find("going")
17
>>> s.isdigit()
False
>>> s.lower()
'so, when is this going to get difficult?'
>>> s.upper()
'SO, WHEN IS THIS GOING TO GET DIFFICULT?'
>>> s.replace("difficult", "easy")
'So, when is this going to get easy?'
>>> s.split()
['So,', 'when', 'is', 'this', 'going', 'to', 'get', 'difficult?']
```



```
>>> s.split("i")
['So, when ', 's th', 's go', 'ng to get d', 'ff', 'cult?']
```

Note the execution of the string method **str.find()** above: `s.find("going")`. This string method returns the first index of the string, “going”, within the string, `s`. Why is the result 17? At first glance, it seems that the first character of the string, “going”, is at position 18. However, strings are sequences (just like lists); therefore, the characters of a string in Python begin at index 0.

### Importing external libraries

It is often useful (and necessary) to import external functionality into our programs. In fact, you've seen (and used) this before (in the lesson *Introduction to Data Structures*), although it may not have been explained in detail. Often, others have designed functions and other bits of code that may prove useful. We don't always want to recreate things that already exist. Python supports the importing of such things via the **import** reserved word. For example, many of the programs we create require the use of mathematical functions beyond simple arithmetic (e.g., sin, cos, tan) or mathematical constants (e.g., pi, e). The structure of an import statement is as follows:

```
import library
```

Pretty simple. Here's an example of the importing and use of the math library:

```
>>> pi

Traceback (most recent call last):
  File "<pyshell#0>", line 1, in <module>
    pi
NameError: name 'pi' is not defined
>>> import math
>>> math.pi
3.141592653589793
>>> math.e
2.718281828459045
>>> math.sin(math.pi)
1.2246467991473532e-16
>>> math.cos(math.pi)
-1.0
>>> math.log(1000)
6.907755278982137
>>> math.log(math.e)
1.0
```

Note in the example, the invalid use of `pi` before importing that math library. In addition, any value or function used in a library must be fully qualified with the name of the library (e.g., we need to specify `math.pi` and not just `pi`). Alternatively, we can itemize what we wish to import from a library. This allows us to use values and functions directly without having to specify the library name. The structure of such an import statement is as follows:

```
from library import function (or constant)
```

For example, the constant PI and the `sin()` function can be formally imported as follows:

```
from math import pi, sin
```

Moreover, these can be directly used as follows:

```
print(pi)
print(sin(pi))
```

Here's the output:

```
3.141592653589793
1.2246467991473532e-16
```

Formally, Python calls its libraries **modules**. And we can even write our own modules! They are just Python programs that typically provide definitions of constants and functions that other Python programs import and make use of. Python modules just need to be saved as a .py file and located in the same folder/directory as a Python program that needs to make use of it. For example, we could include several useful functions in a file called MyGoodies.py. Suppose that it contained the following:

```
from time import time

# starts a timer
def start_time():
    global start
    start = time()
# stops the timer and returns the time elapsed
def stop_time():
    stop = time()
    elapsed = stop - start

    return elapsed
```

The purpose of this example module is to use it to time how long algorithms take to execute. It's quite simple. The function *start\_time* effectively starts a timer (through the time library's *time* function) by capturing the current “time” – which is essentially the number of seconds elapsed since an epoch defined in your operating system. For Unix and Unix-like operating systems (e.g., the “flavor” of Linux used on the Raspberry Pi), the epoch is 1970-01-01 00:00:00. The function *stop\_time* captures the current time again (this time, after the algorithm has finished), and calculates and returns the difference between the two.

We can make use of this module as follows:

```
from MyGoodies import start_time, stop_time

# start a timer
start_time()

# do something that takes a little time
for i in range(100000000):
    pass

# stop the timer
```

```
duration = stop_time()

# display how long it took
print("Algorithm took {} seconds.".format(duration))
```

Note that the “algorithm” in the test code above really does nothing. It's just there to take up some noticeable amount of time so that the module can be tested. Here's a sample run:

```
jgourd@pi:~$ python MyGoodiesTest.py
Algorithm took 5.71634602547 seconds.
```